

# GPT2PPO: Auto Regressive Proximal Policy Optimization

1<sup>st</sup> Andrei Cozma

Department of Electrical Engineering & Computer Science  
University of Tennessee  
Knoxville, United States  
acozma@vols.utk.edu

2<sup>nd</sup> Hunter Price

Department of Electrical Engineering & Computer Science  
University of Tennessee  
Knoxville, United States  
hprice7@vols.utk.edu

**Abstract**—This document is a model and instructions for L<sup>A</sup>T<sub>E</sub>X. This and the IEEEtran.cls file define the components of your paper [title, text, heads, etc.]. \*CRITICAL: Do Not Use Symbols, Special Characters, Footnotes, or Math in Paper Title or Abstract.

**Index Terms**—Reinforcement Learning, PPO, GPT2

## I. INTRODUCTION

In this project we explore the use of transformers in the context of Reinforcement Learning. The majority of theoretical works assume that problems follow a Markovian process, which is not always the case. Some problems need the context of previous states and actions to make an informed decision on the next decision. As a result, we propose an addition to the basic Proximal Policy Optimization (PPO) algorithm by using the Generative Pre-trained Transformer 2 (GPT2) model as the encoder for the critic network. This will allow the critic network to take into account the context of previous states and actions as well as apply attention to past states and actions that may be important. We will test this model on the LunarLander-v2 and Acrobot-v1 OpenAI Gym environments with discrete action spaces and compare it to the original PPO algorithm. Additionally we will test the model on BipedalWalker-v3 with continuous action spaces and compare it to the original PPO algorithm.

## II. PREVIOUS WORK

## III. BACKGROUND

All of the environments used in this project are from the OpenAI Gym library [1]. The environments are described below.

### A. Lunar Lander

The Lunar Lander environment is a rocket trajectory optimization problem<sup>1</sup> shown in Figure 1. The goal is to actuate the lander to the landing pad at coordinates (0,0) without crashing. OpenAI Gym offers two versions of the environment: discrete or continuous. In this work we only use the discrete version. The state space is a 8-dimensional vector containing the x and y positional coordinates of the agent, its x and

y linear velocities, its angle, its angular velocity, and two booleans that represent whether each leg is in contact with the ground or not. The action space is a single discrete scalar with values ranging from 0 to 3. The values correspond to the following actions: do nothing, fire left orientation engine, fire main engine, fire right orientation engine. The reward structure contains both positive and negative rewards. If the lander moves away from the landing pad, it gains a negative reward. If the lander crashes, it receives an -100 reward. If it comes to rest, it receives an +100 reward. Each leg with ground contact is +10 points. Firing the main engine is -0.3 points each frame. Firing the side engine is -0.03 points each frame. The lander's initial state is at the top center of the environment with a random initial force applied to its center. The episode ends if the lander crashes, goes outside of the viewport, or comes to a resting position.



Fig. 1. The Lunar Lander environment.

### B. Acrobot

Open AI Gym's implementation of the Acrobot environment<sup>2</sup> is based off the work of Sutton and Barto [2] shown

<sup>1</sup>OpenAI Gym Lunar Lander: [https://www.gymnasium.dev/environments/box2d/lunar\\_lander](https://www.gymnasium.dev/environments/box2d/lunar_lander)

<sup>2</sup>OpenAI Gym Acrobot: [https://www.gymnasium.dev/environments/classic\\_control/acrobot](https://www.gymnasium.dev/environments/classic_control/acrobot)

in Figure 2. The environment is a 2-link pendulum with only the second joint actuated with a discrete action space. The goal is to swing the end of the pendulum up to a given height. The state space is a 6-dimensional vector containing the sin and cos of the two joint angles and the joint angular velocities. The action space is a single discrete scalar with values ranging from 0 to 5. The values correspond to the following actions: apply +1, 0, or -1 torque to the actuated joint. The reward structure contains only negative rewards. At each timestep the agent receives a reward of -1 for each step that does not reach the goal. If the goal is reached the agent receives a reward of 0. The episode ends if the agent reaches the goal height or if the episode exceeds the maximum number of timesteps.

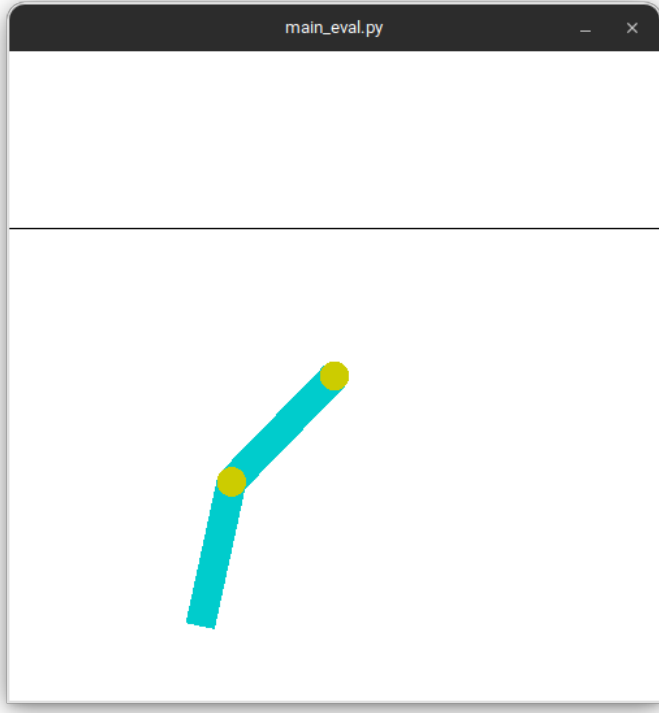


Fig. 2. The Acrobot environment.

### C. Bipedal Walker

The Bipedal Walker environment is a 2D simulation of a bipedal walker robot with 4-joints<sup>3</sup> shown in Figure 3. The goal is to keep the walker upright for as long as possible. The state space is a 24-dimensional vector containing: the hull angle speed, angular velocity, horizontal speed, vertical speed, position of joints and joints angular speed, legs contact with ground, and 10 lidar rangefinder measurements. Notably, there are no coordinates given in the state vector. The action space is a 4 dimensional vector containing continuous values of each joints motor speed between -1 and 1. The reward structure contains both positive and negative rewards. A positive reward

is given for moving forward. If the agent falls it receives a reward of -100. Applying motor torque costs a small amount of reward. The agent's initial state is standing at the left of the terrain with the hull horizontal, with both legs in the same position with a slight knee angle. The episode terminates if the agent's hull makes contact with the ground or if the agent reaches the end of the terrain length.



Fig. 3. The Bipedal Walker environment.

## IV. METHODOLOGY

### A. Reinforcement Learning Methods

### V. CODE DESIGN

#### A. Utilized Libraries

In this work we chose to use pytorch as our main deep learning library. Pytorch is a popular deep learning library that is well documented and has a large community. To help maintain the structure of the models we also used the Pytorch-Lightning library. Pytorch-Lightning is a high-level library that allows for easy training and testing of models. Pytorch-Lightning also provides a module named Lightning-Bolts that contains many pre-built models and utilities. For the GPT2 implementation we used the HuggingFace Transformers library. We use the OpenAI Gym Library to provide the environments for our models to train and test on [1]. Finally, we use wandb to log the results of all of our experiments [3].

#### B. File Structure

This works file structure is as follows. The main directory contains the **main\_train.py** and **main\_eval.py** files. These files are used to train and test the models respectively. The core code is contained within the **rlib** directory. Within the **rlib/examples** contains an A2C and PPO baseline implementations that were taken from the Lightning-Bolts library, copied into our project.

Within the **rlib** directory there are many files. The important files are the files that are primarily used are: **Model.py**, **GPT2PPO.py**, **CommonGPT2.py**, **CommonBase.py**, and

<sup>3</sup>OpenAI Gym Acrobot: <https://www.gymnasium.dev/environments/box2d/bipedal-walker>

**GPT2.py**. These files will be discussed in more detail in the following sections.

### C. Training and Testing Tools

To train a model we use the **main\_train.py** file. This file will train a model on a given environment and log the results. To train a model we use the following command:

```
python3 main_train.py \
    -m ppo_gpt \
    -e LunarLander-v2
```

For the **main\_train.py** file we have the following additional command line arguments:

- `-e` or `--env` - The Open AI Gym environment [Default: LunarLander-v2]
- `-m` or `--model_name` - The name of the model to train [Default: `ppo_ex` (PPO)]
- `-ne` or `--num_epochs` - The number of epochs to train for [Default: 150]
- `--al_check_interval` - The interval of epochs to run validation [Default: 5]
- `--wandb_project` - The wandb project to log to [Default: `rl_project`]
- `--wandb_entity` - The wandb entity to log to [Default: `ece517`]
- `--seed` - The seed value for training [Default: 123]

The training command will train the model then save the trained model to the checkpoints directory with the following structure:

```
checkpoints/
<model_name>/
  <env_name>/
    <model_hash>.pt
```

To test a model (watch it interact in the environment) we use the **main\_eval.py**. This file will load a model from the checkpoints directory and run it in the environment. To test a model we use the following command:

```
python3 main_eval.py -f <model_checkpoint>
```

For the **main\_eval.py** file we have the following additional command line arguments:

- `-f` or `--file_path` - The path to the model checkpoint
- `-ne` or `--num_episodes` - The number of episodes to run [Default: 5]
- `--running_rew_len` - The length of rewards to store at once [Default: 50]
- `--seed` - The seed value for testing [Default: 123]

Both the **main\_train.py** and **main\_eval.py** files utilize the **rllib/Model.py** file. This file abstracts the training and evaluation process from the training and testing commands into a class named `Model`. It simply loads in the correct model and trains or tests it.

### D. Models

The baseline PPO model we use to compare against our results is in the **rllib/examples/PPOExample.py** file. In it contains code taken from the Lightning-Bolts library with no modifications. We store this model locally to protect against any future api changes in the library. The file contains a class named `PPO` that inherits from the `LightningModule` class. Our implementation is built upon this class.

This works implementation resides at **rllib/GPT2PPO.py**. This file contains the `GPT2PPO` class which implements the same interface as the reference `PPO`.

## VI. RESULTS

## VII. CONCLUSION

## REFERENCES

- [1] G. Brockman, V. Cheung, L. Pettersson, J. Schneider, J. Schulman, J. Tang, and W. Zaremba, "Openai gym," *arXiv preprint arXiv:1606.01540*, 2016.
- [2] R. S. Sutton and A. G. Barto, *Reinforcement learning: An introduction*. MIT press, 2018.
- [3] L. Biewald, "Experiment tracking with weights and biases," 2020, software available from wandb.com. [Online]. Available: <https://www.wandb.com/>

## APPENDIX

### A. Team Member Contributions

**Andrei Cozma**. text here.

**Hunter Price**. text here.